

# Retrieve Roads from Aerial Imagery Using Deep Learning

Capstone Project Mid-Point Presentation, CPLN 680, Spring 2022, UPenn

Jiamin Tan

# Object

**Learn how to train a neural network to detect roads from aerial images.**

# Object

**Learn how to train a neural network to detect roads from aerial images.**

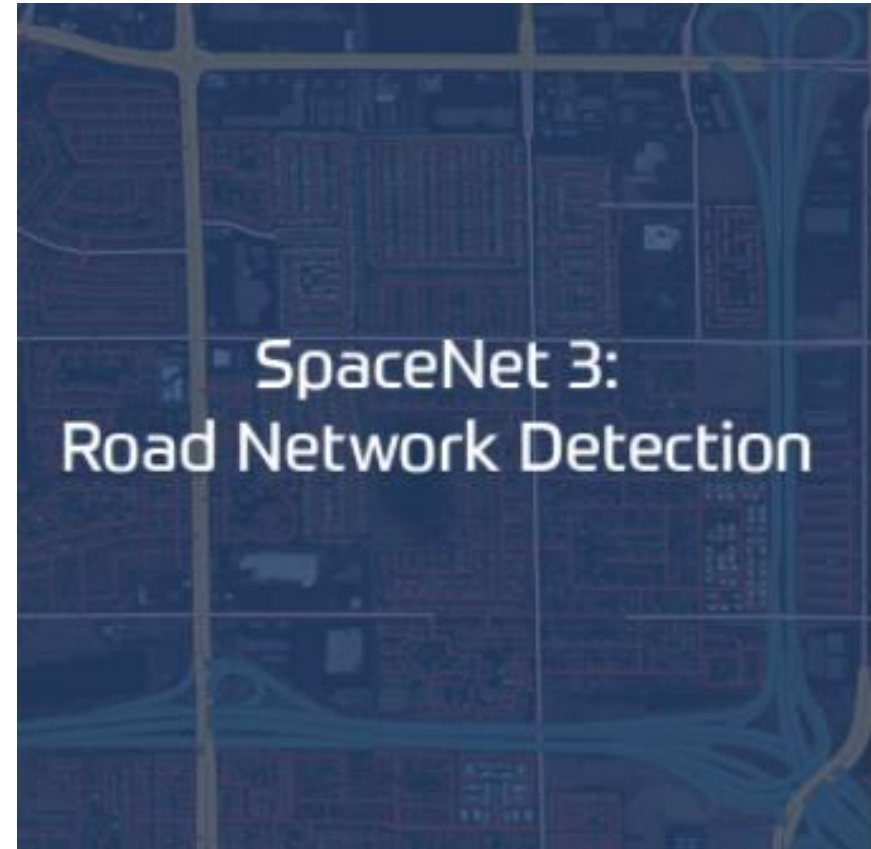
**Current Status:**

I am learning 😊

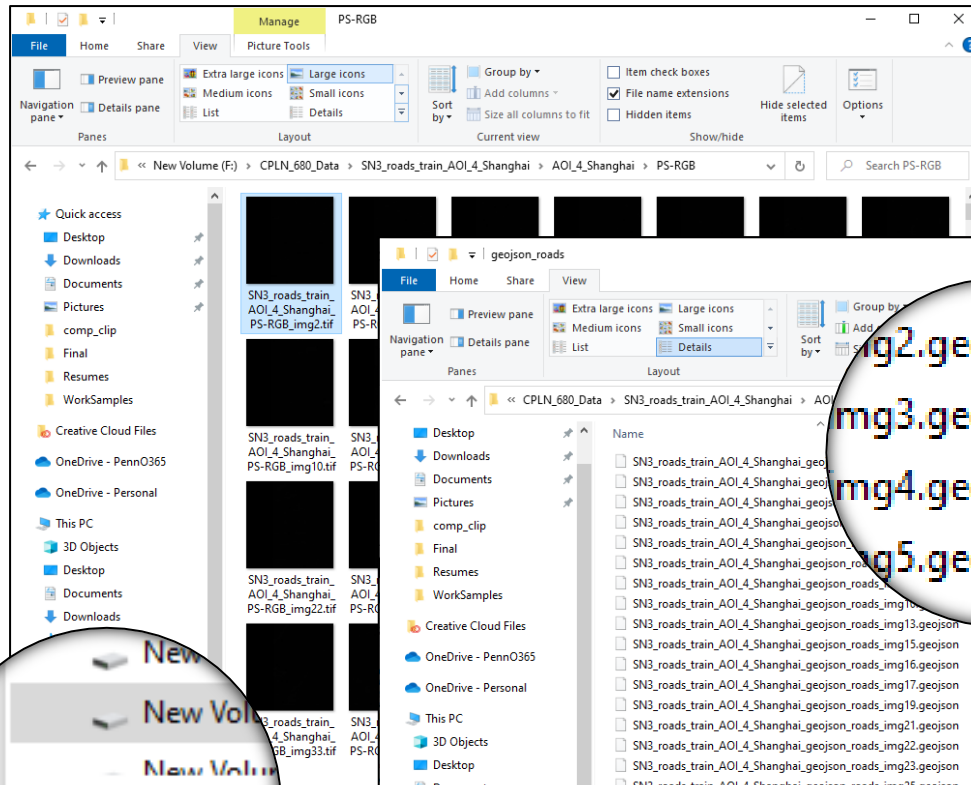
but my algorithm is not learning 😞

which means both of us haven't learnt enough so far...

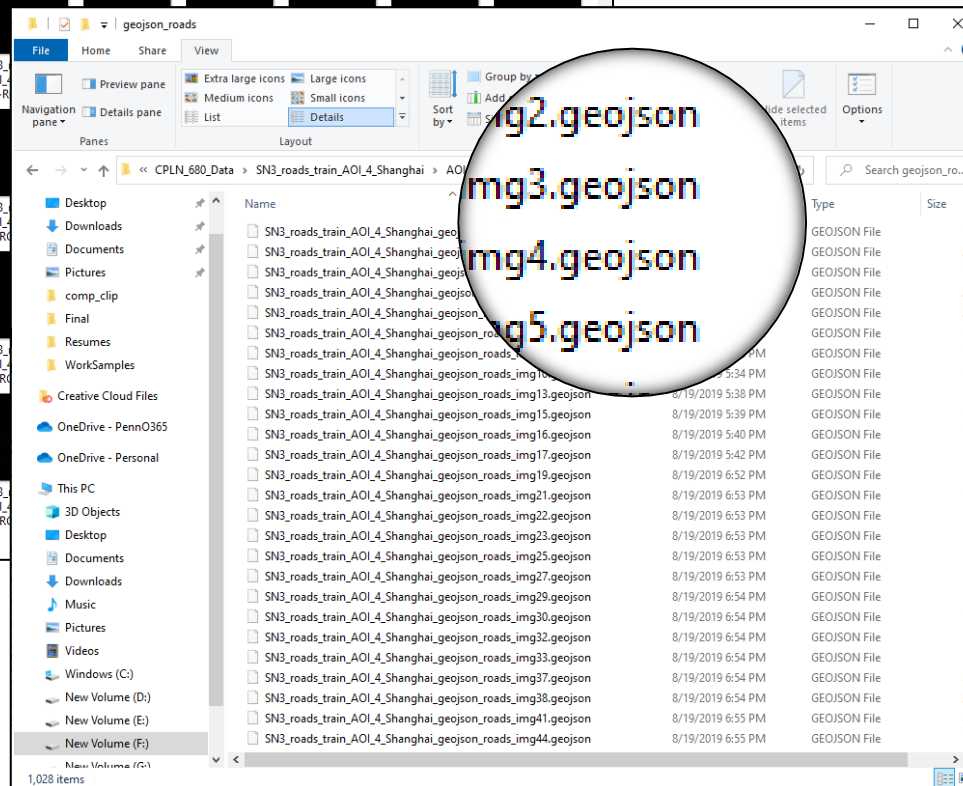
# Main Data Source



## Training Image Set



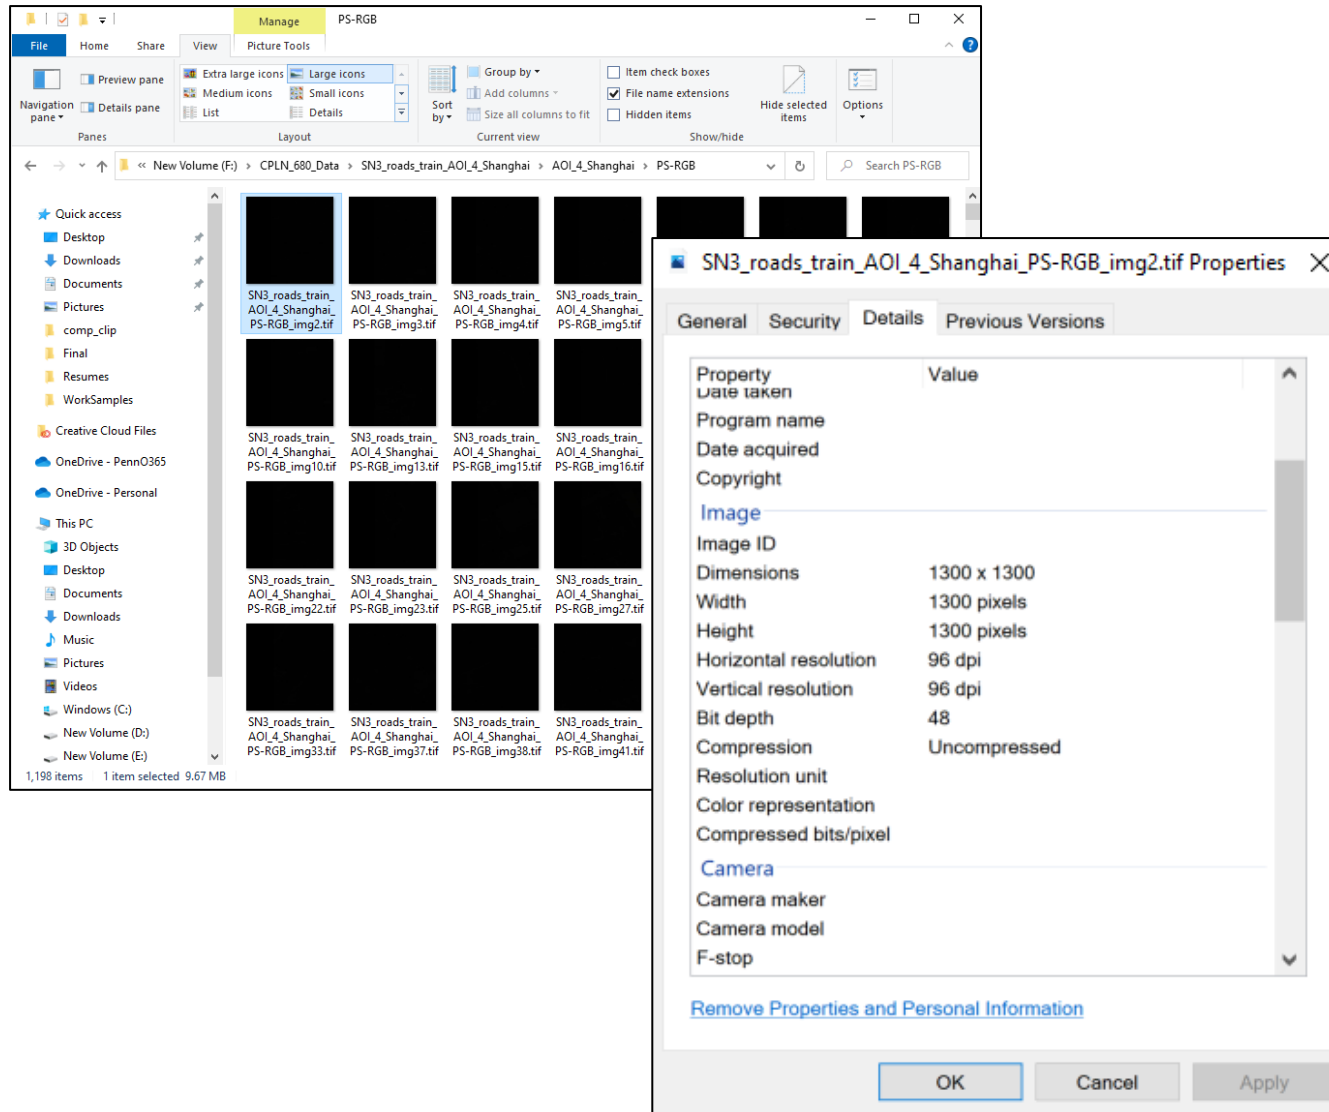
## Labels



.tif files stored satellite images.

.geojson files stored strings (roads) as ground truth.

## Training Image Set



For each pan-sharpened Image:

1300 x 1300 pixels.

Each pixel has a spatial resolution of 0.31m x 0.31m.

Each tile is, therefore, 400m x 400m.

Bit depth is 48, so 16 bit for each band, and the value of a pixel in each band is from 0 to 65535.





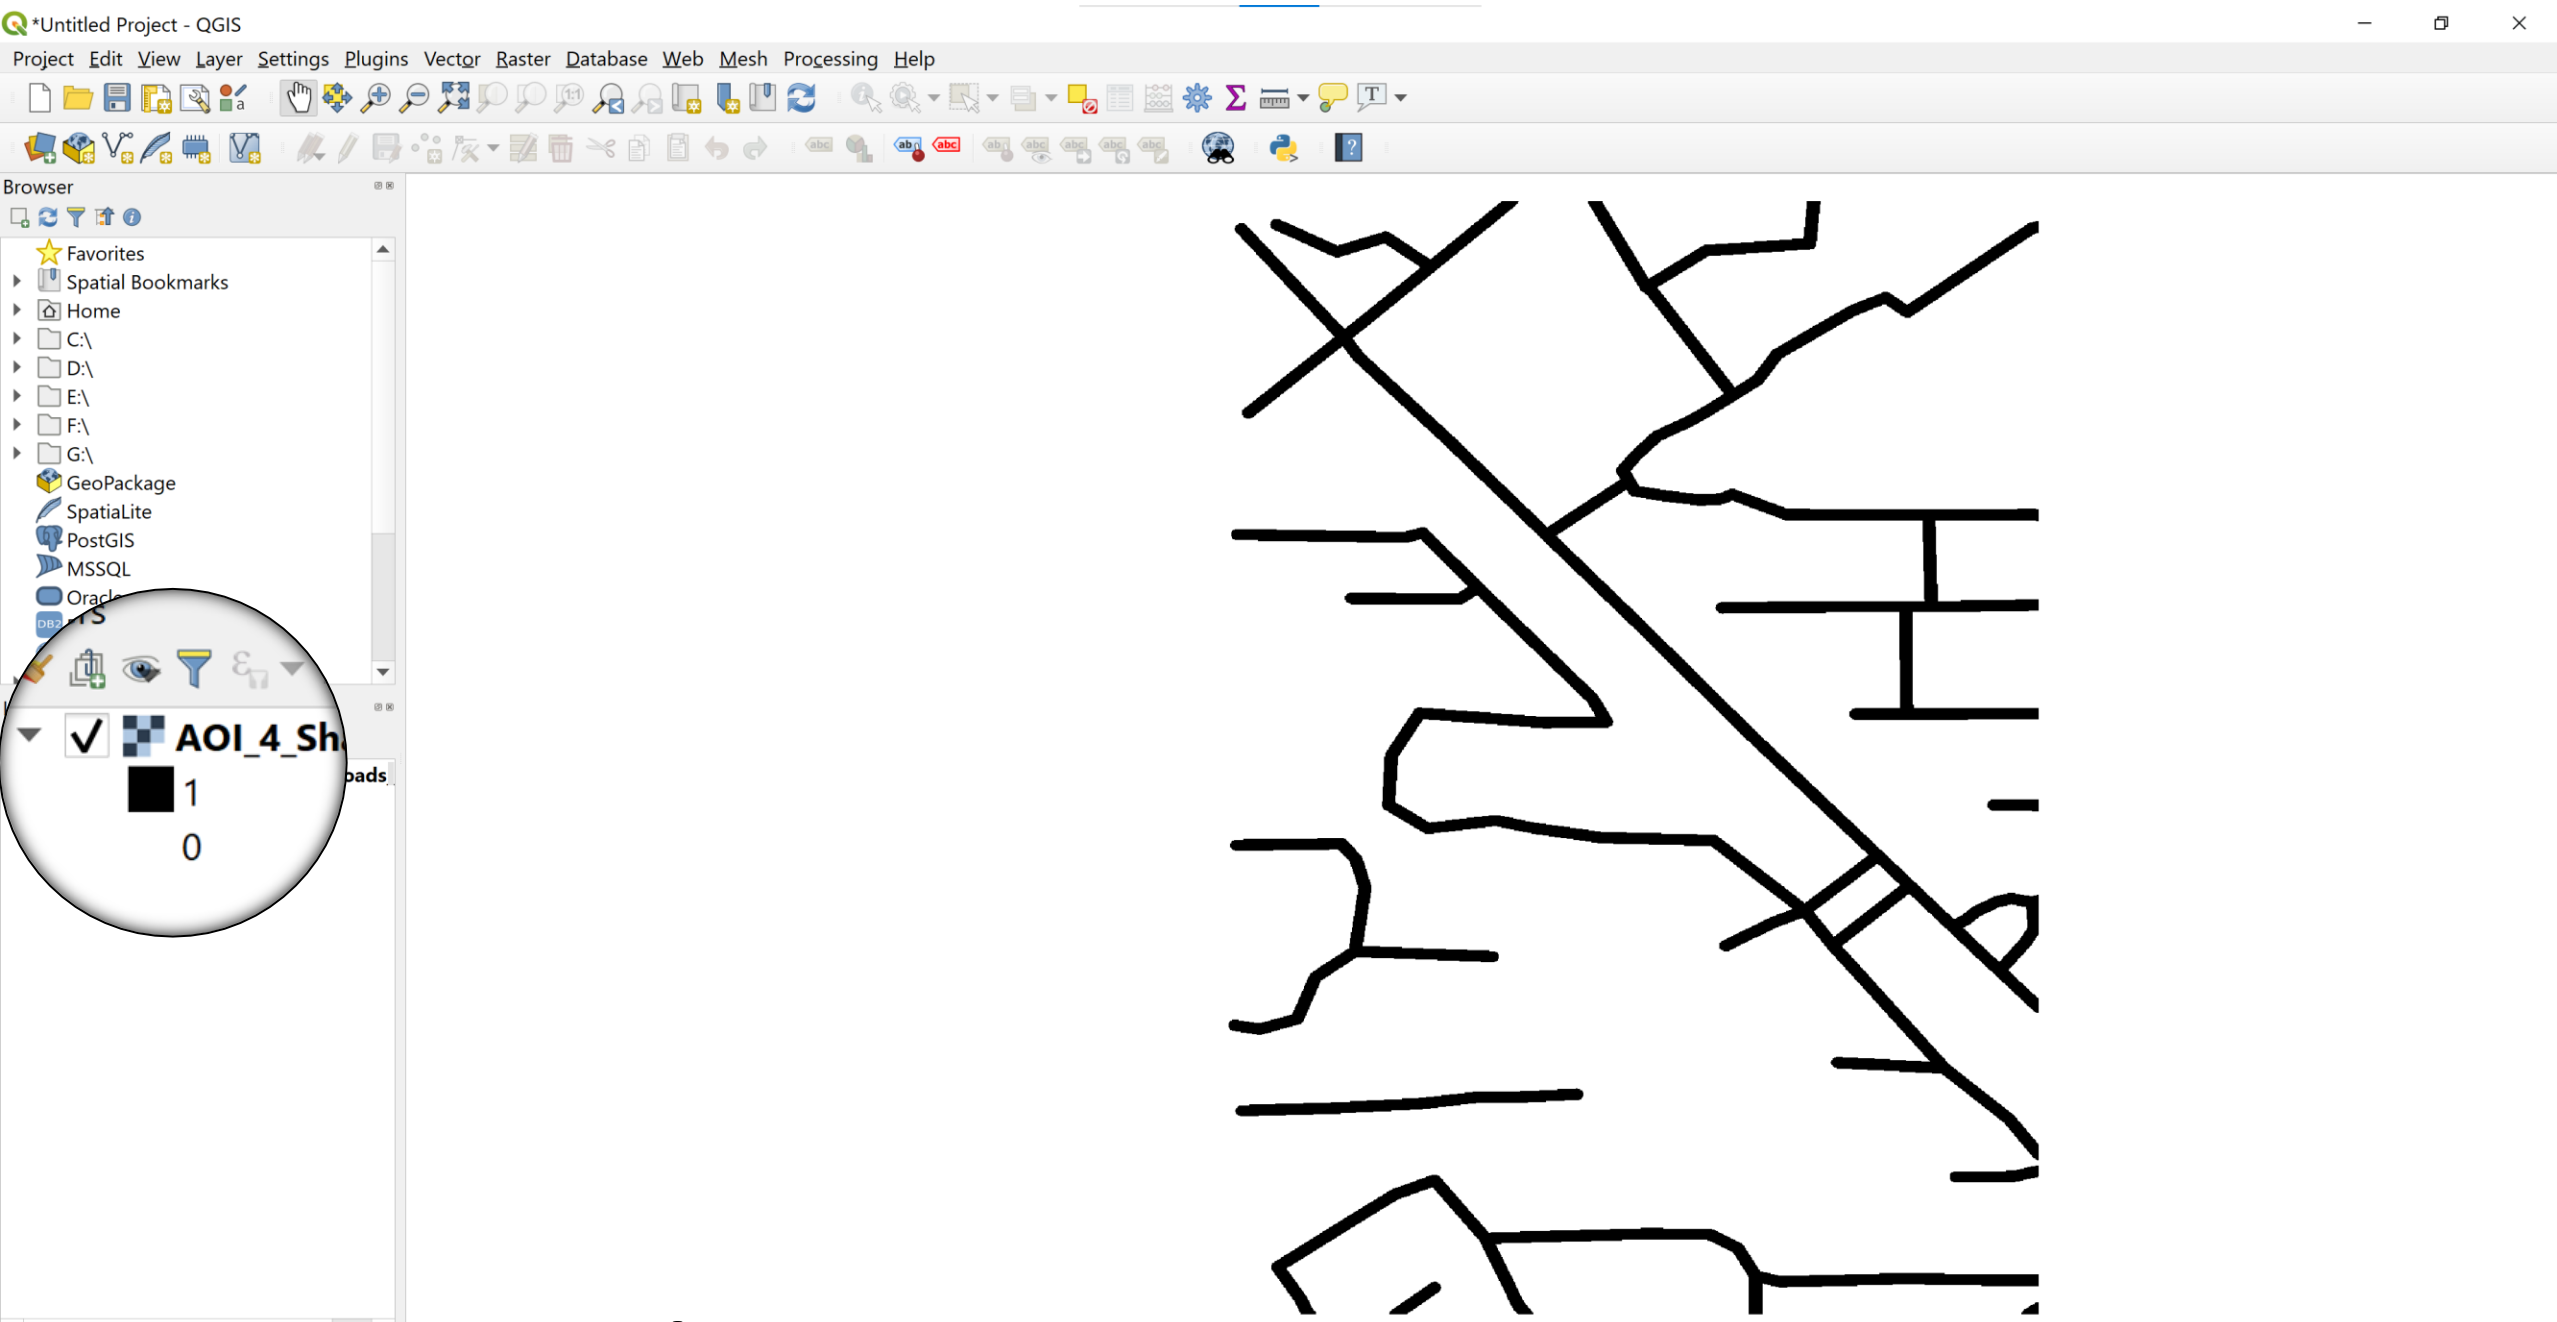
Browser

- ★ Favorites
- ▶ Spatial Bookmarks
- Home
- C:\
- D:\
- E:\
- F:\
- G:\
- GeoPackage
- SpatiaLite
- PostGIS
- MSSQL
- Oracle
- DB2

Legend

- ☒ SN3 roads
- ☒ SN3\_roads

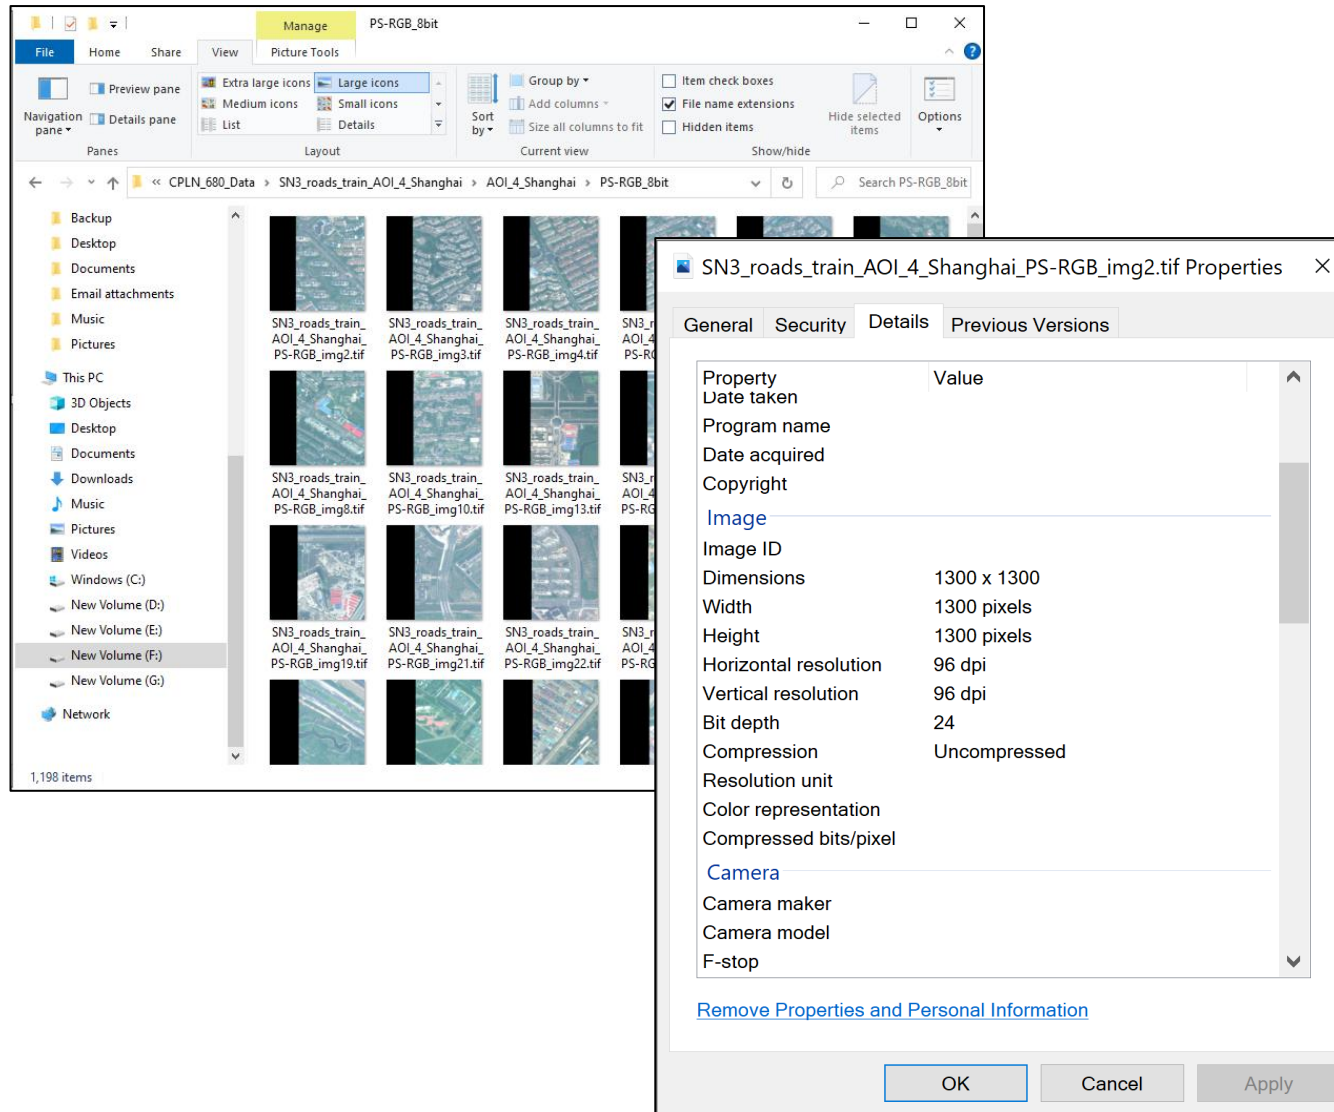




Create Binary Masks



## Training Image Set (8-bit)



Each Image **Now**:

1300 x 1300 pixels.

Each pixel has a spatial resolution of 0.31m x 0.31m.

Each tile is, therefore, 400m x 400m.

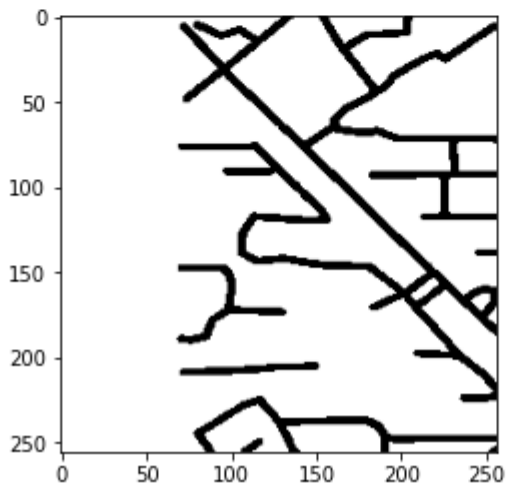
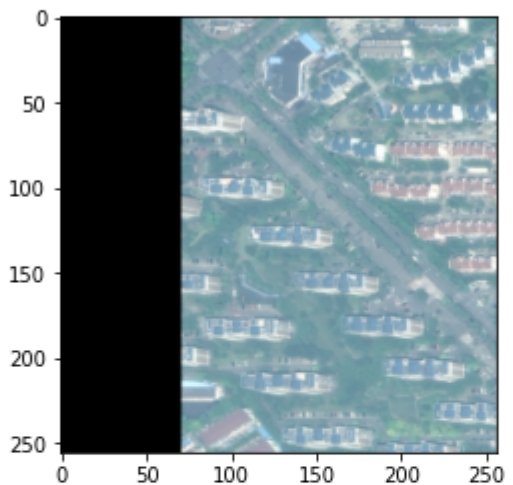
Bit depth is **24**, so **8** bit for each band – value of a pixel in each band is from **0** to **255**.

# Create 8-Bit Images

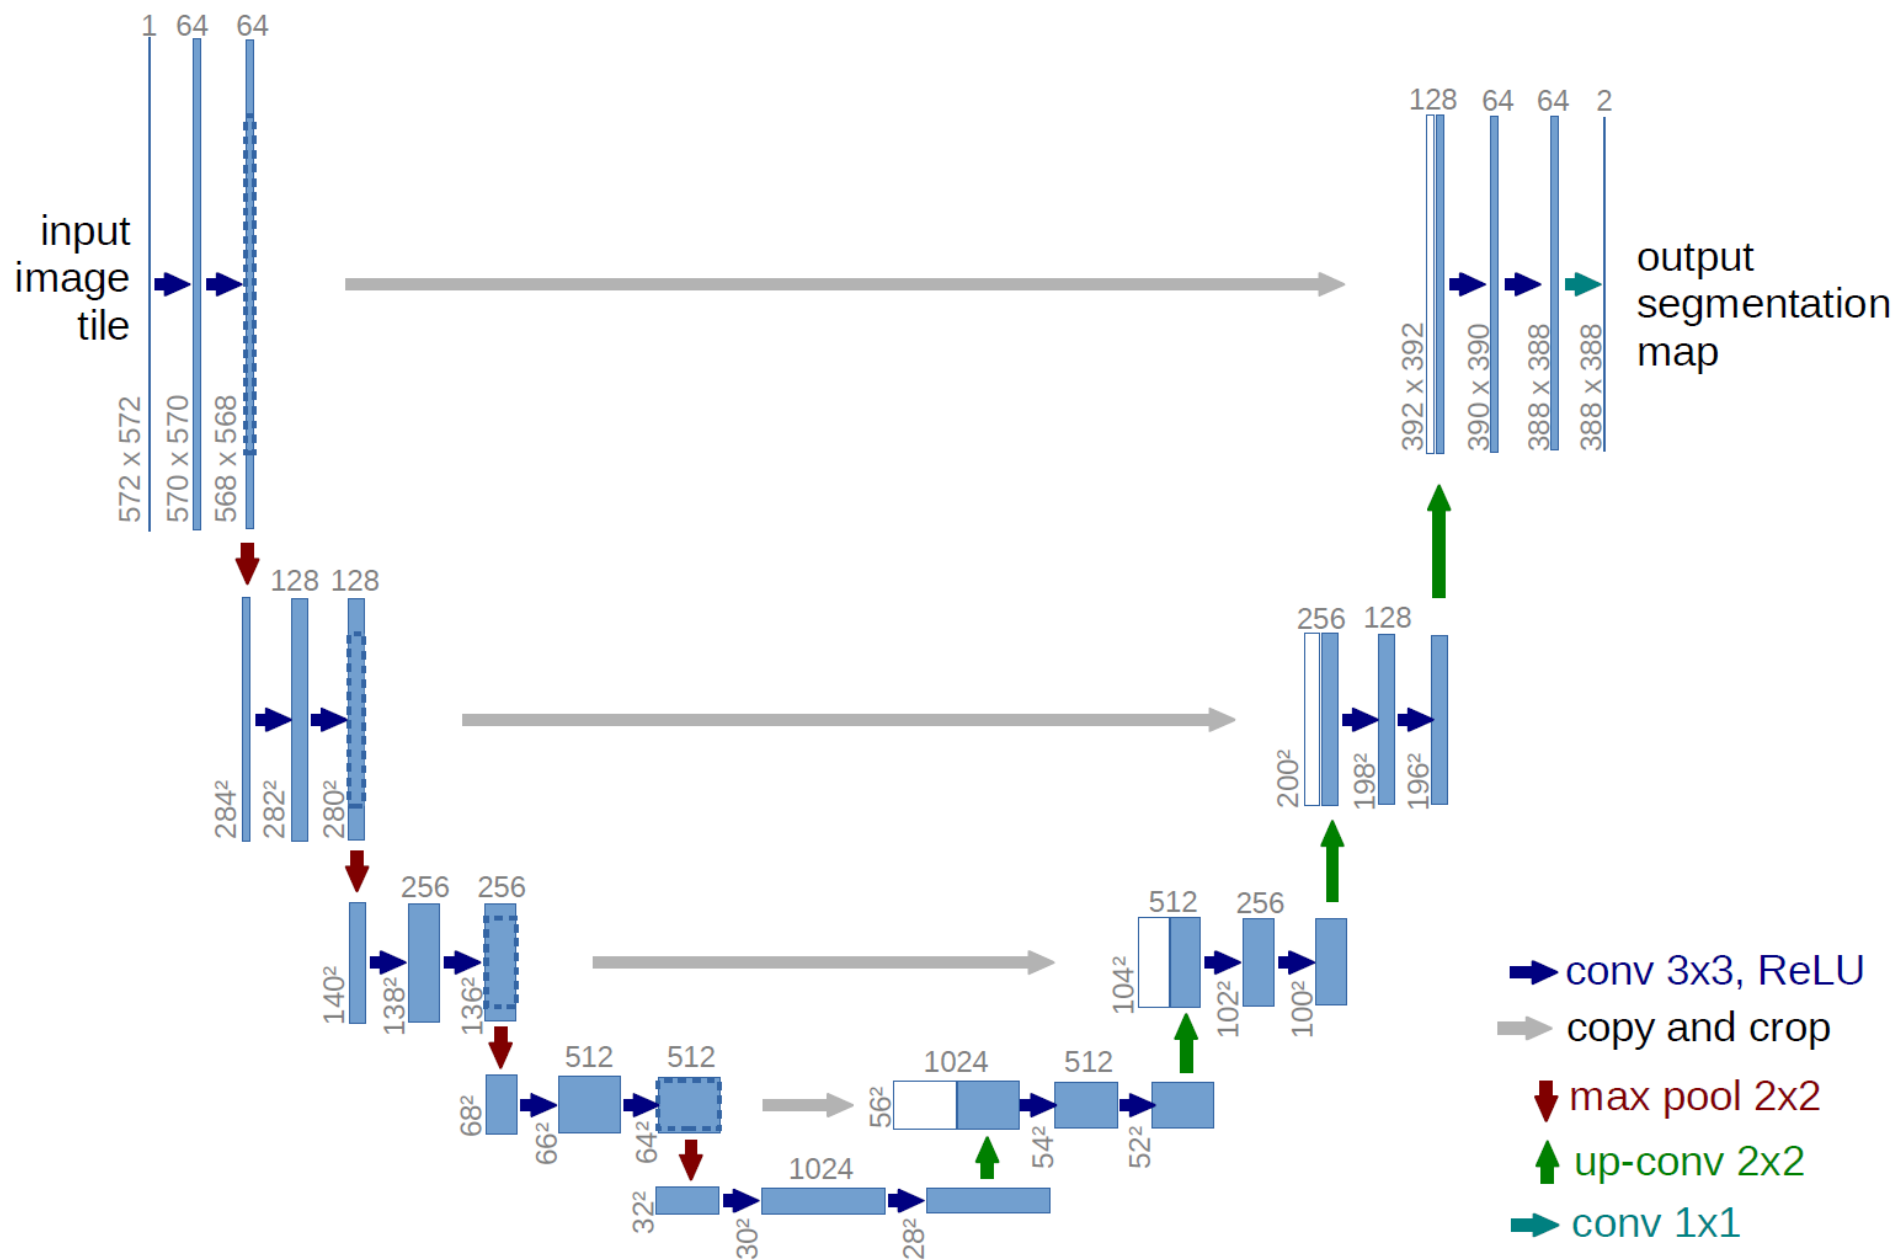
✓  
0s



```
image_x = random.randint(0, len(train_img_ids))  
plt.imshow(X_train[train_img_ids.index('SN3_roads_train_AOI_4_Shanghai_PS-RGB_img2.tif')])  
plt.show()  
plt.imshow(np.squeeze(Y_train[train_img_ids.index('SN3_roads_train_AOI_4_Shanghai_PS-RGB_img2.tif')]), cmap='Greys')  
plt.show()
```



# Load Training Set into Python



# U-Net by Ronneberger et al. (2015)

```
[19] c1 = tf.keras.layers.Conv2D(64, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(inputs)
c1 = tf.keras.layers.Dropout(0.1)(c1)
c1 = tf.keras.layers.Conv2D(64, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c1)
p1 = tf.keras.layers.MaxPooling2D((2,2))(c1)

[20] c2 = tf.keras.layers.Conv2D(128, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(p1)
c2 = tf.keras.layers.Dropout(0.1)(c2)
c2 = tf.keras.layers.Conv2D(128, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c2)
p2 = tf.keras.layers.MaxPooling2D((2,2))(c2)

[21] c3 = tf.keras.layers.Conv2D(256, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(p2)
c3 = tf.keras.layers.Dropout(0.1)(c3)
c3 = tf.keras.layers.Conv2D(256, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c3)
p3 = tf.keras.layers.MaxPooling2D((2,2))(c3)

[22] c4 = tf.keras.layers.Conv2D(512, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(p3)
c4 = tf.keras.layers.Dropout(0.1)(c4)
c4 = tf.keras.layers.Conv2D(512, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c4)
p4 = tf.keras.layers.MaxPooling2D((2,2))(c4)

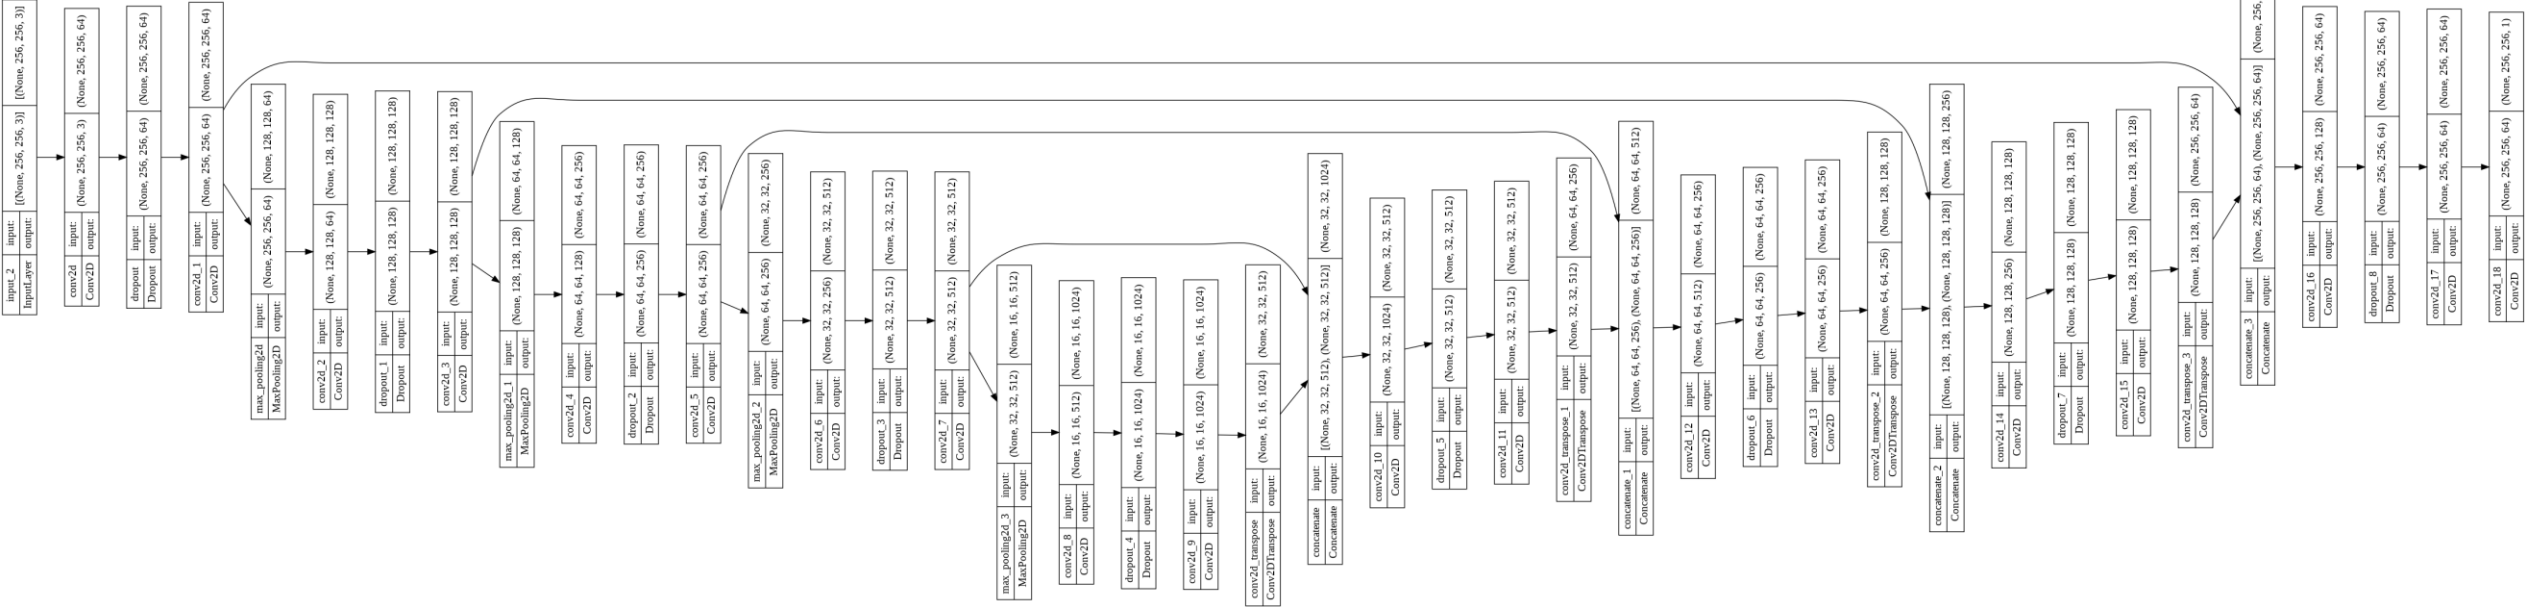
[23] c5 = tf.keras.layers.Conv2D(1024, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(p4)
c5 = tf.keras.layers.Dropout(0.3)(c5)
c5 = tf.keras.layers.Conv2D(1024, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c5)

[24] u6 = tf.keras.layers.Convolution2DTranspose(512, (2,2), strides = (2,2), padding = 'same')(c5)
u6 = tf.keras.layers.concatenate([u6, c4])
c6 = tf.keras.layers.Conv2D(512, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(u6)
c6 = tf.keras.layers.Dropout(0.2)(c6)
c6 = tf.keras.layers.Conv2D(512, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c6)

[25] u7 = tf.keras.layers.Convolution2DTranspose(256, (2,2), strides = (2,2), padding = 'same')(c6)
u7 = tf.keras.layers.concatenate([u7, c3])
c7 = tf.keras.layers.Conv2D(256, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(u7)
c7 = tf.keras.layers.Dropout(0.2)(c7)
c7 = tf.keras.layers.Conv2D(256, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(c7)

[26] u8 = tf.keras.layers.Convolution2DTranspose(128, (2,2), strides = (2,2), padding = 'same')(c7)
u8 = tf.keras.layers.concatenate([u8, c2])
c8 = tf.keras.layers.Conv2D(128, (16,1), activation = 'relu', kernel_initializer = 'he_normal', padding = 'same')(u8)
```

# Hard Code U-Net Using Keras

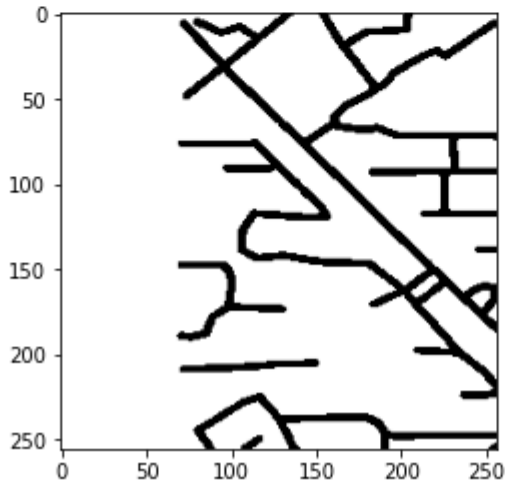
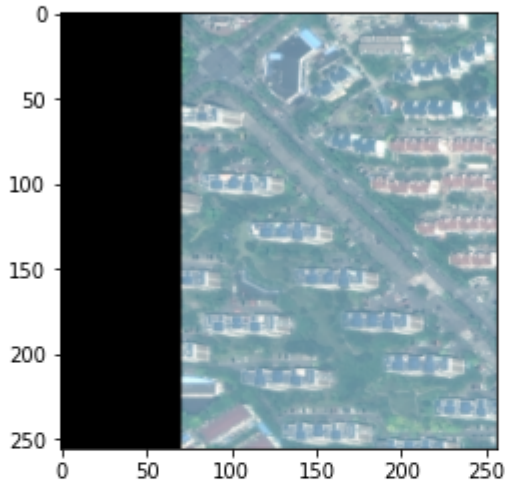


# Hard Code U-Net Using Keras (Cont.)

✓  
0s



```
image_x = random.randint(0, len(train_img_ids))  
plt.imshow(X_train[train_img_ids.index('SN3_roads_train_AOI_4_Shanghai_PS-RGB_img2.tif')])  
plt.show()  
plt.imshow(np.squeeze(Y_train[train_img_ids.index('SN3_roads_train_AOI_4_Shanghai_PS-RGB_img2.tif')]), cmap='Greys')  
plt.show()
```



Each Image **After Resizing**:

**256 x 256** pixels.

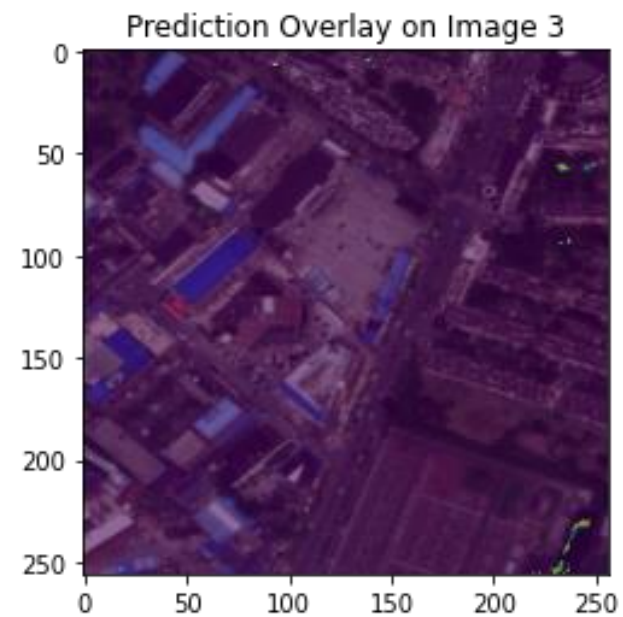
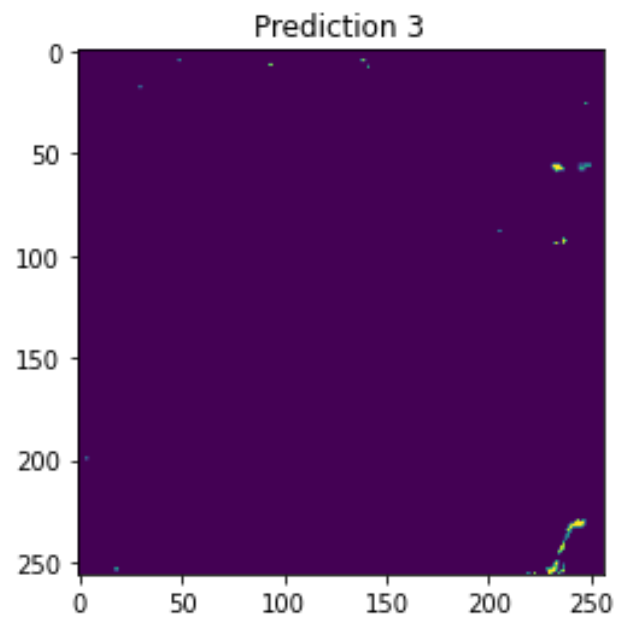
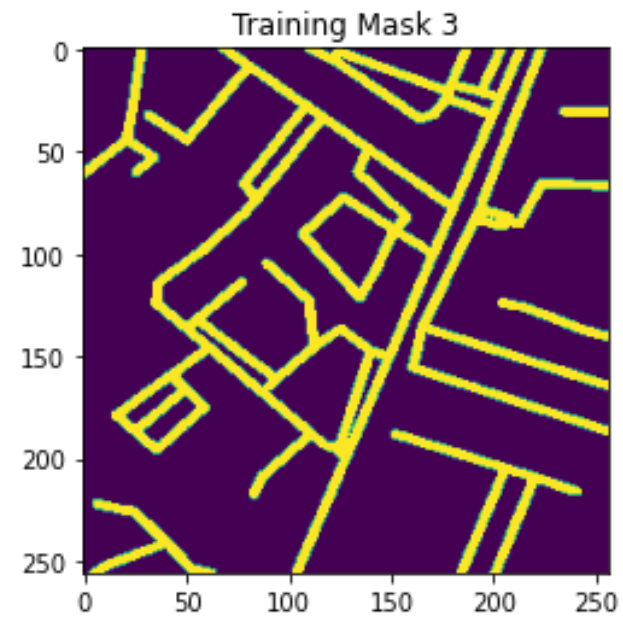
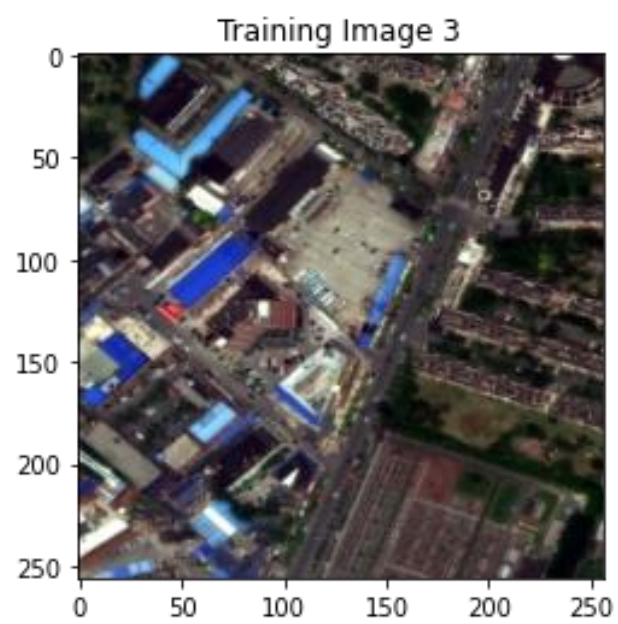
Each pixel has a spatial resolution of  
**1.56m x 1.56m**.

Each tile is, still, 400m x 400m.

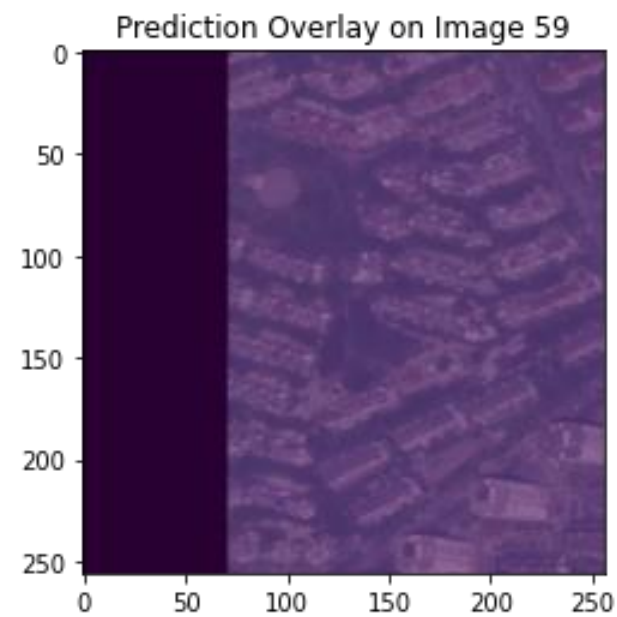
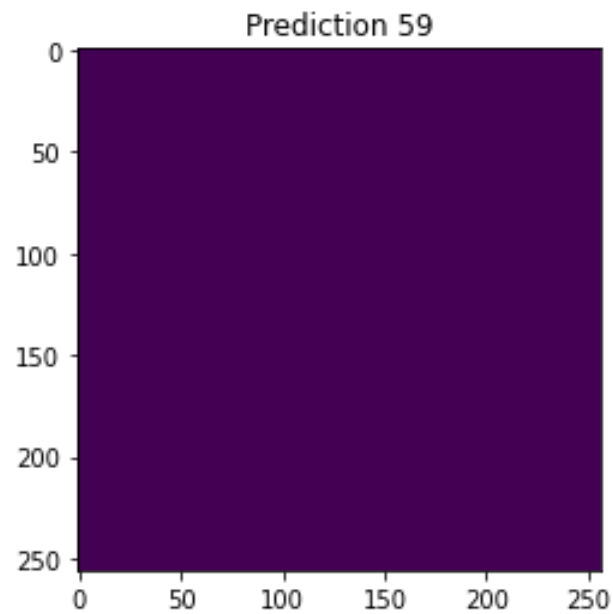
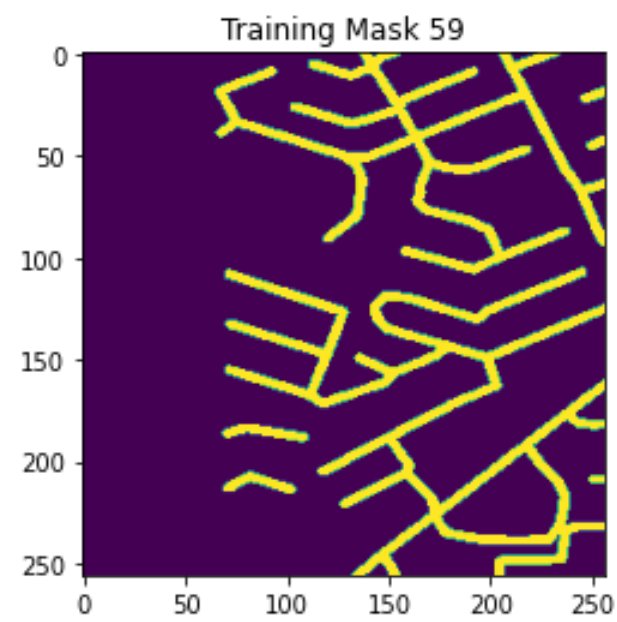
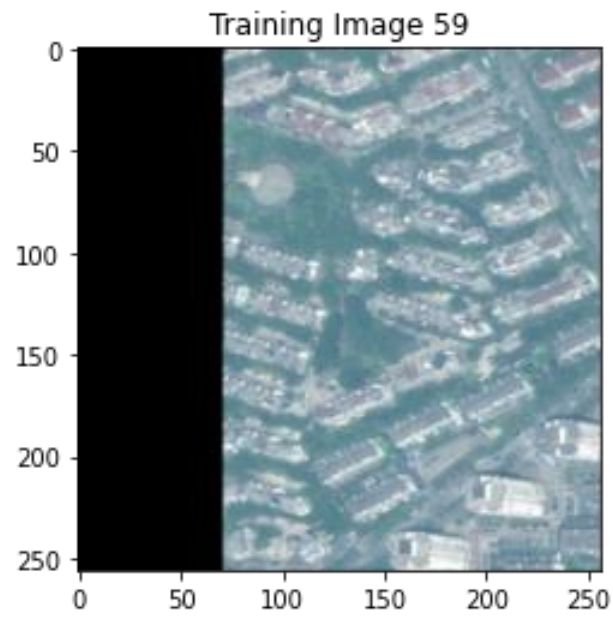
Bit depth is **24**, so **8** bit for each band –  
value of a pixel in each band is from **0** to **255**.

**Train 100 images with reduced size (256, 256)**

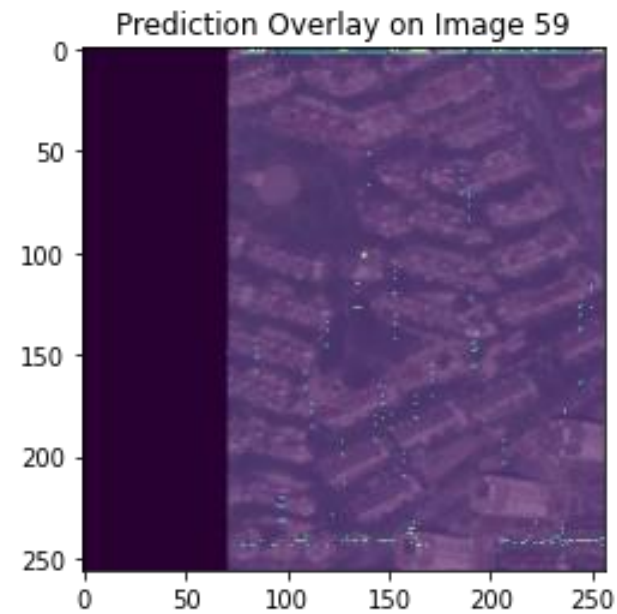
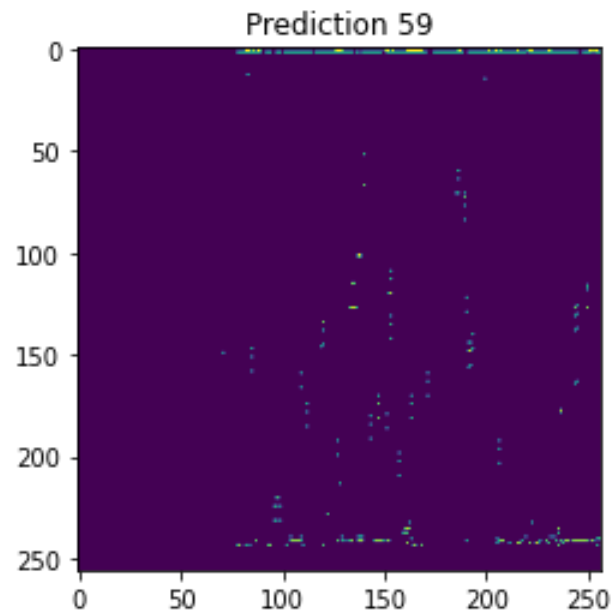
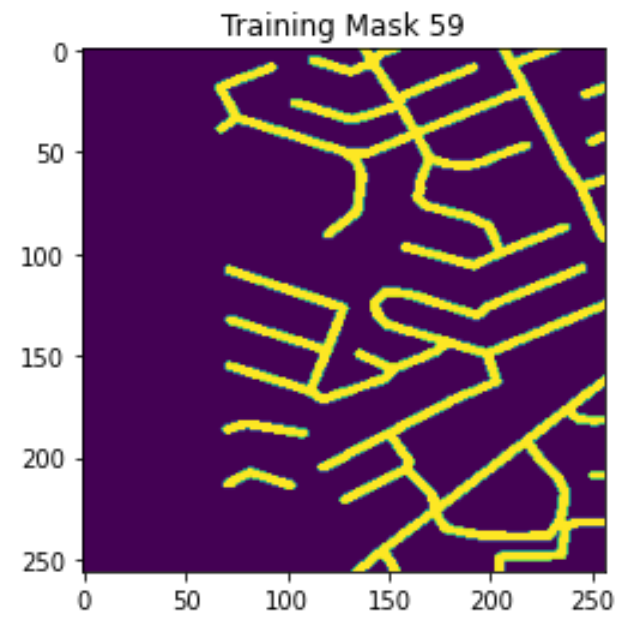
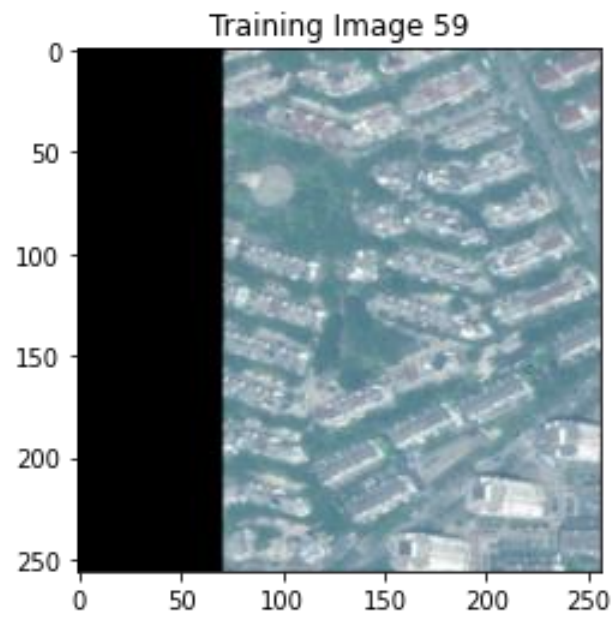




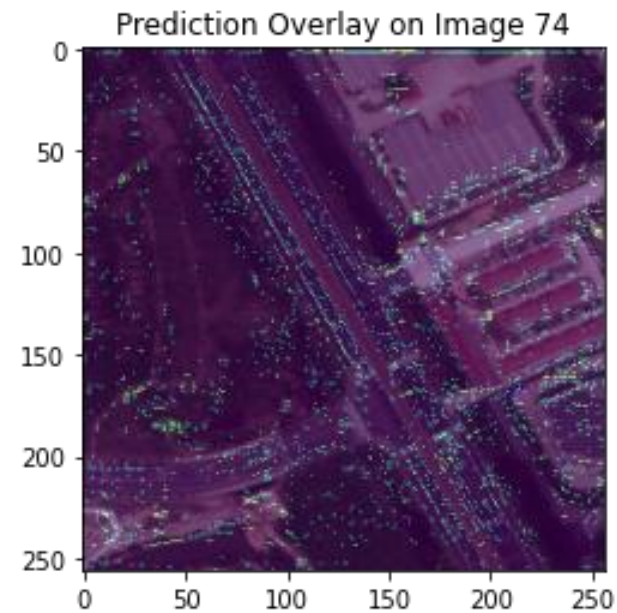
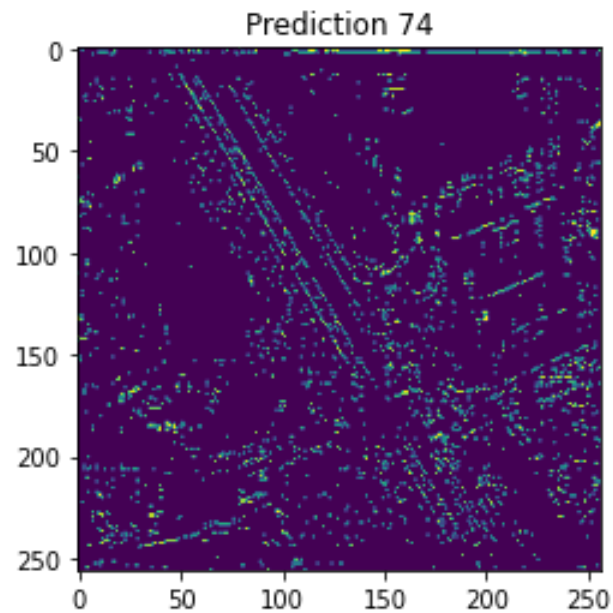
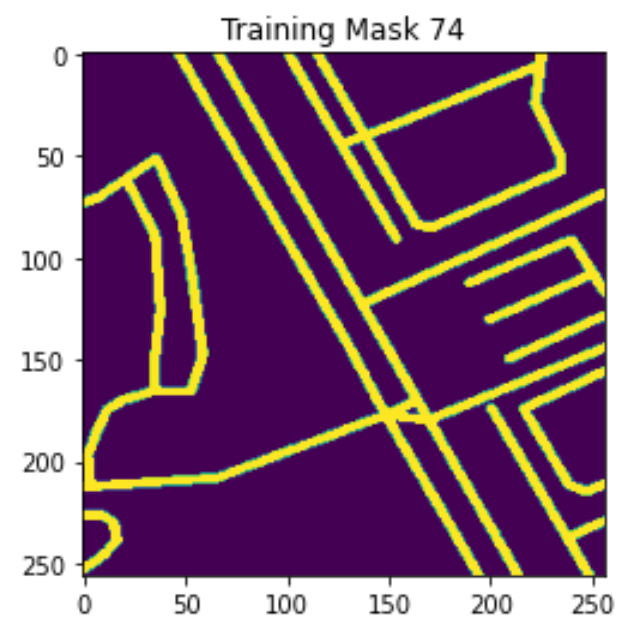
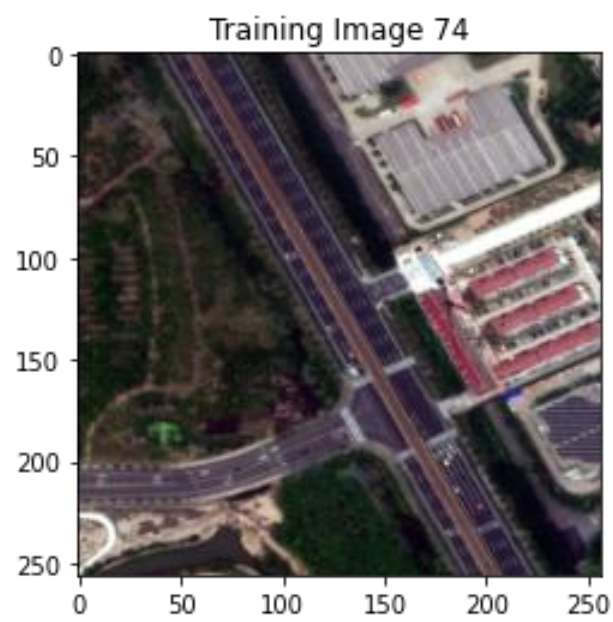
**The results are not good (threshold = 0.5)...**



**The results are not good (threshold = 0.5)...**

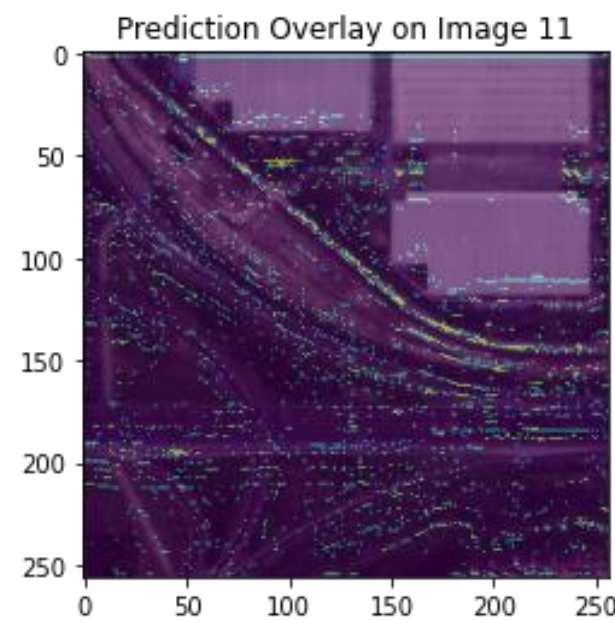
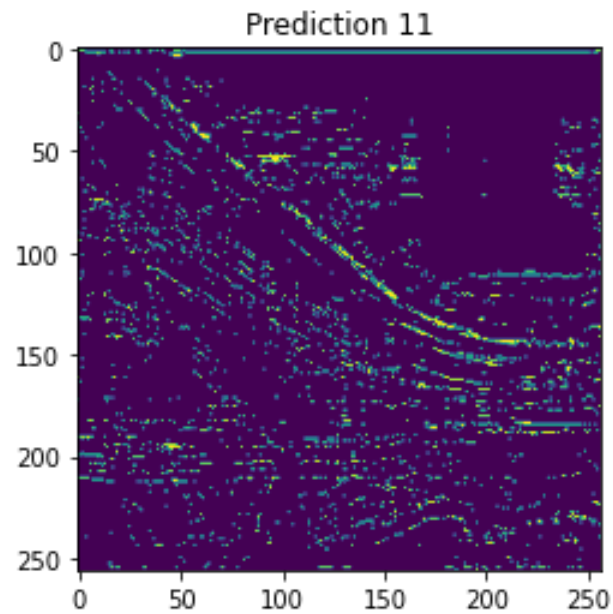
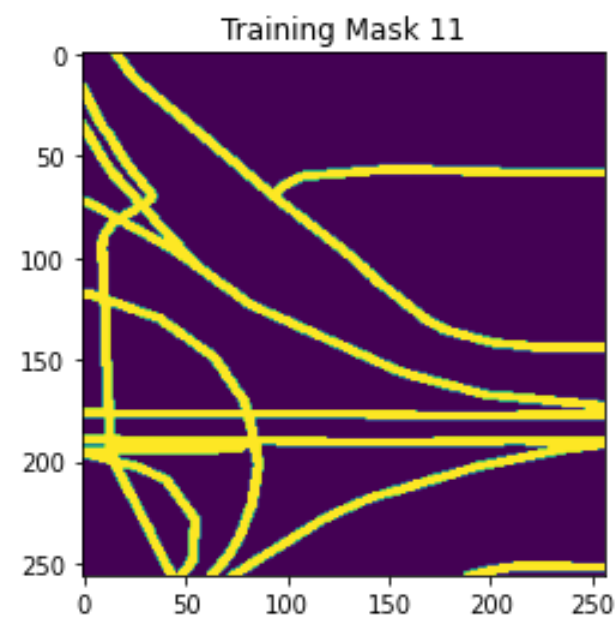
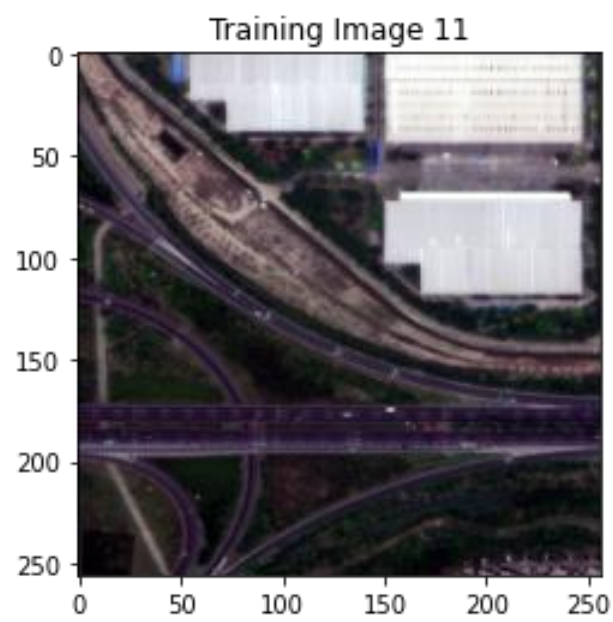


**The same training image in another run (threshold = 0.5)...**

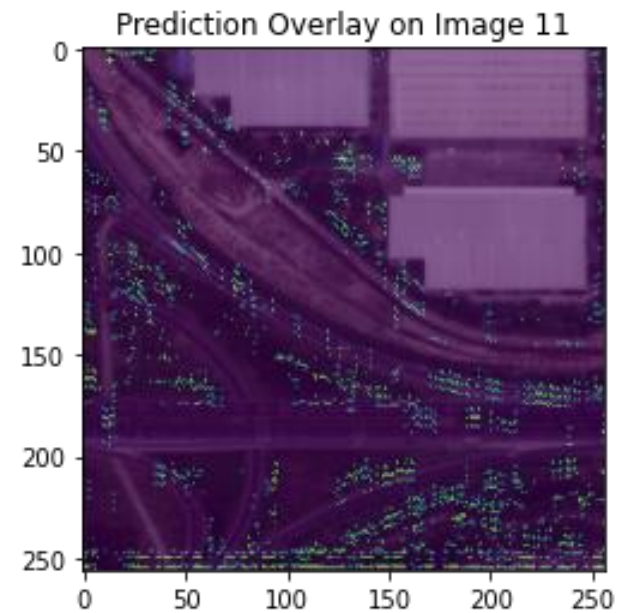
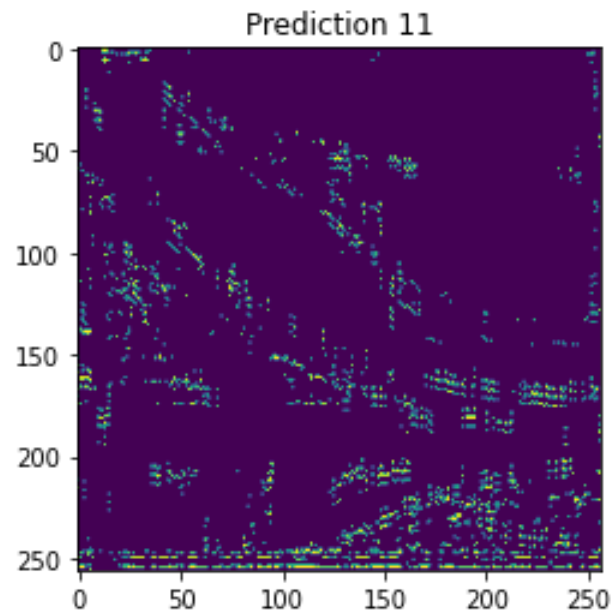
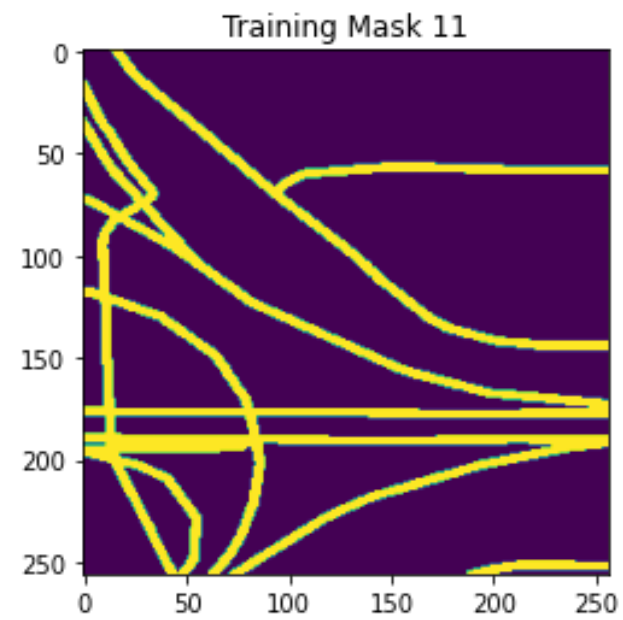
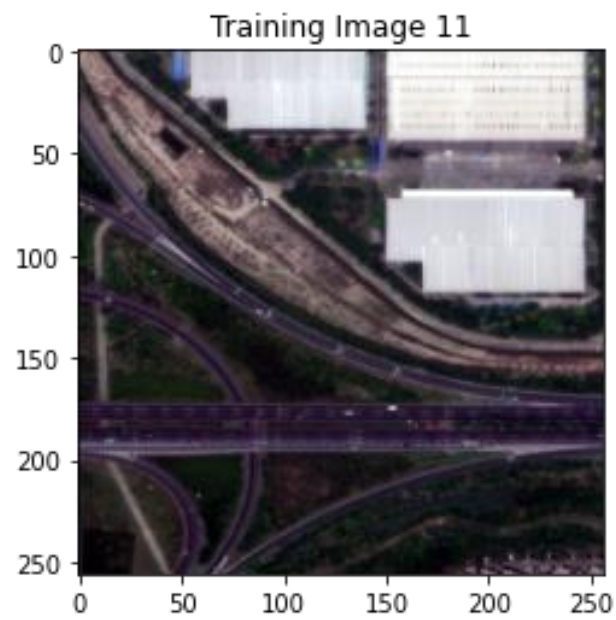


**Sometimes it sort of learnt (threshold = 0.5)...**



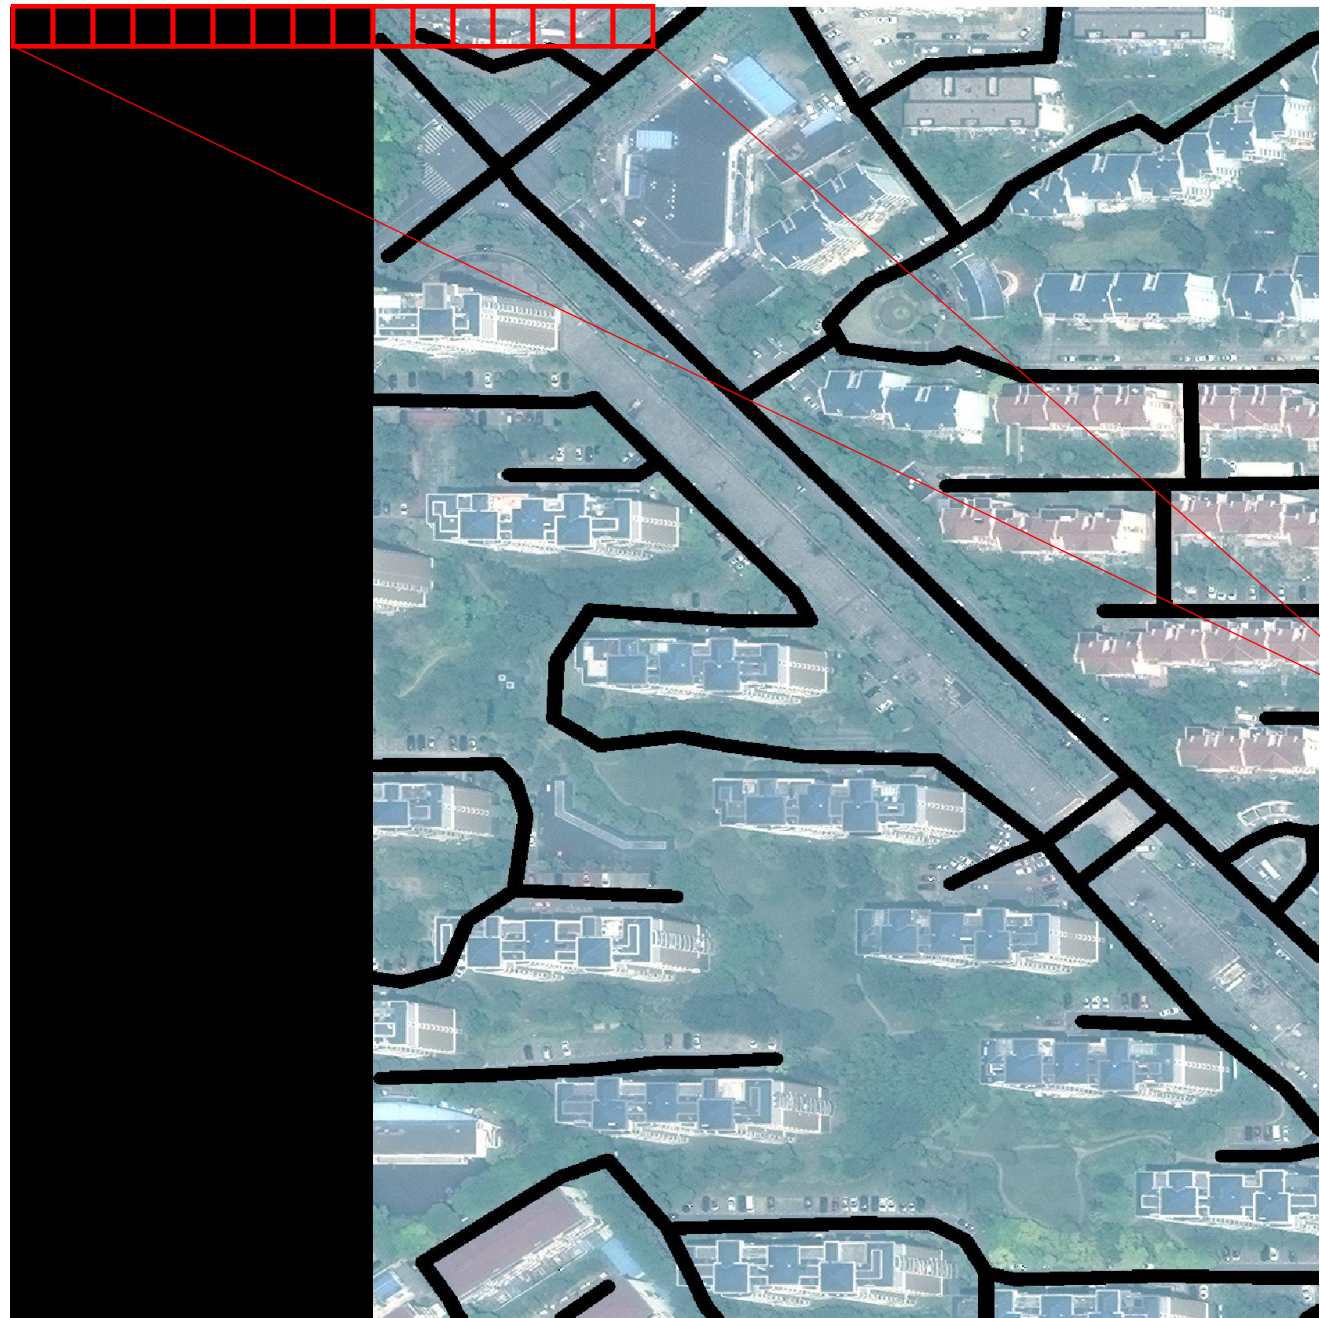


**Sometimes it sort of learnt (threshold = 0.5)...**

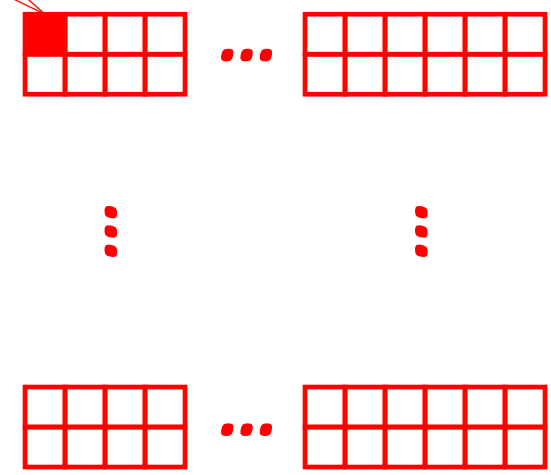


**Or it learnt the wrong thing in another run (threshold = 0.3)...**



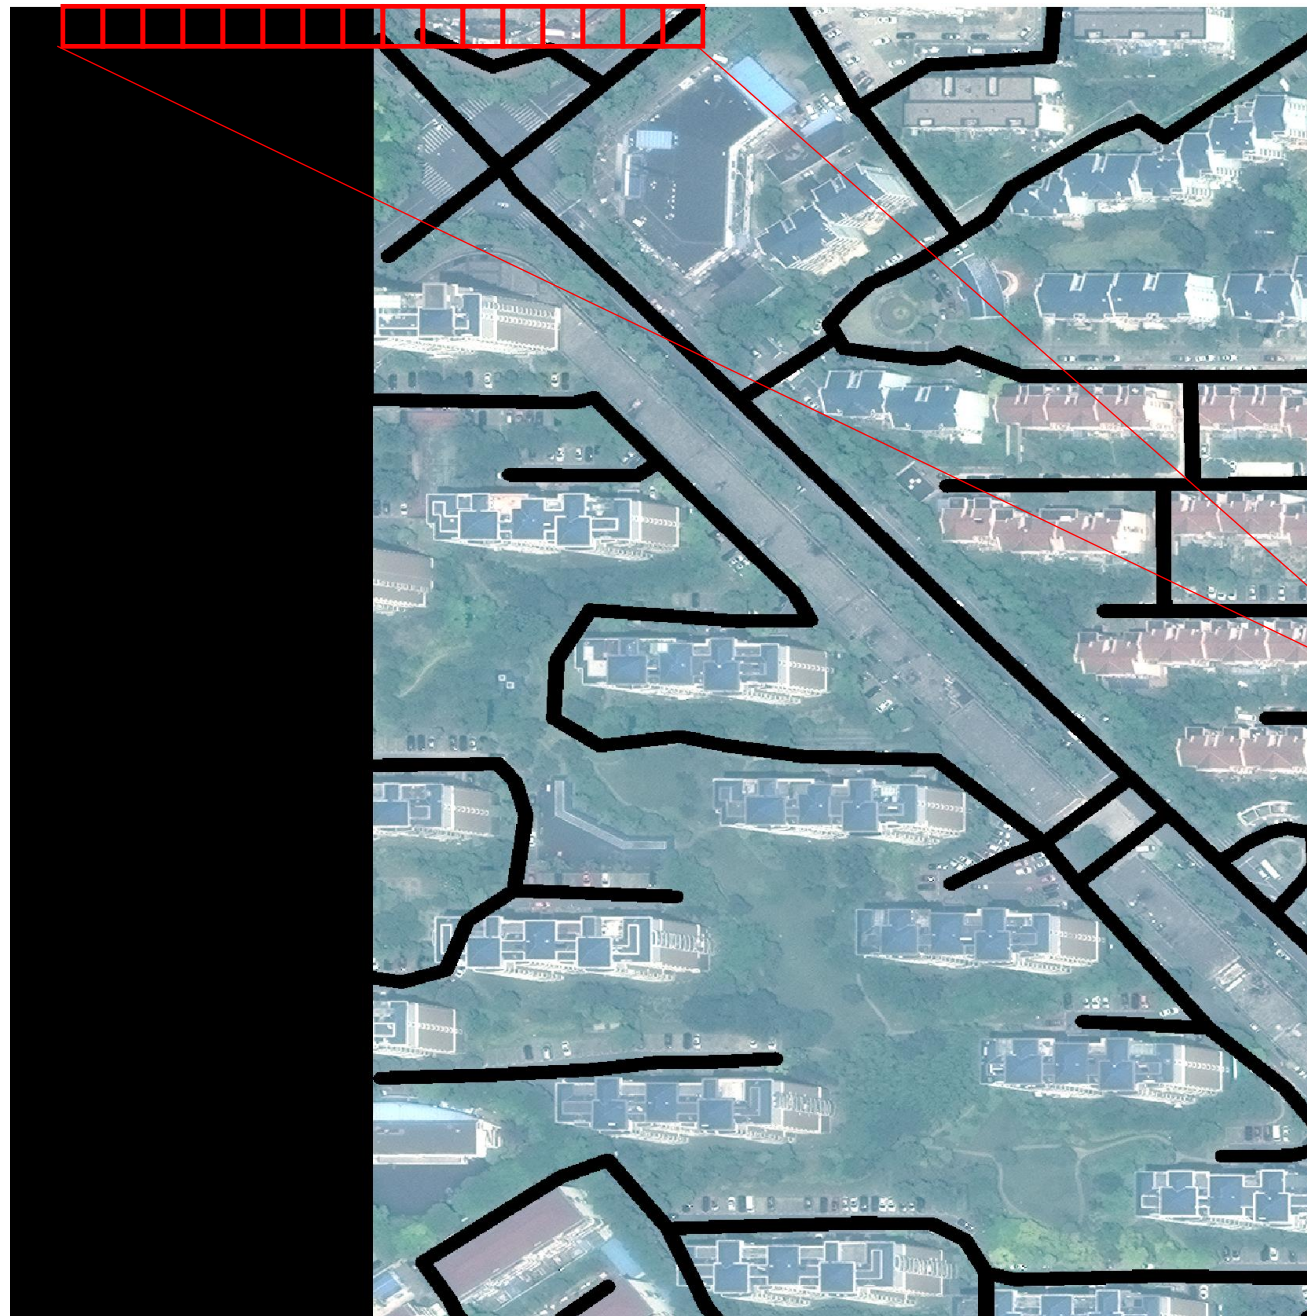


Hidden Layer

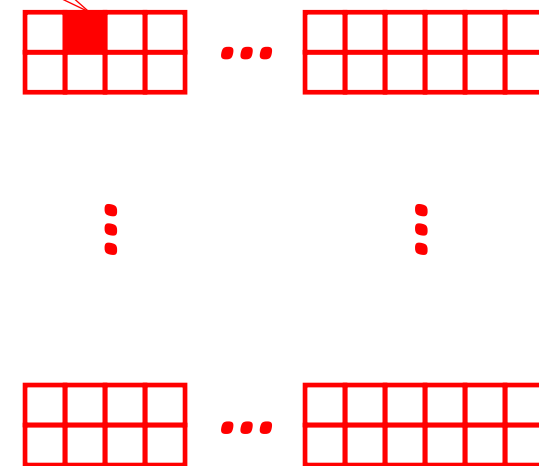


A (16, 1) instead of (3, 3) Kernel (*not in real scale*)



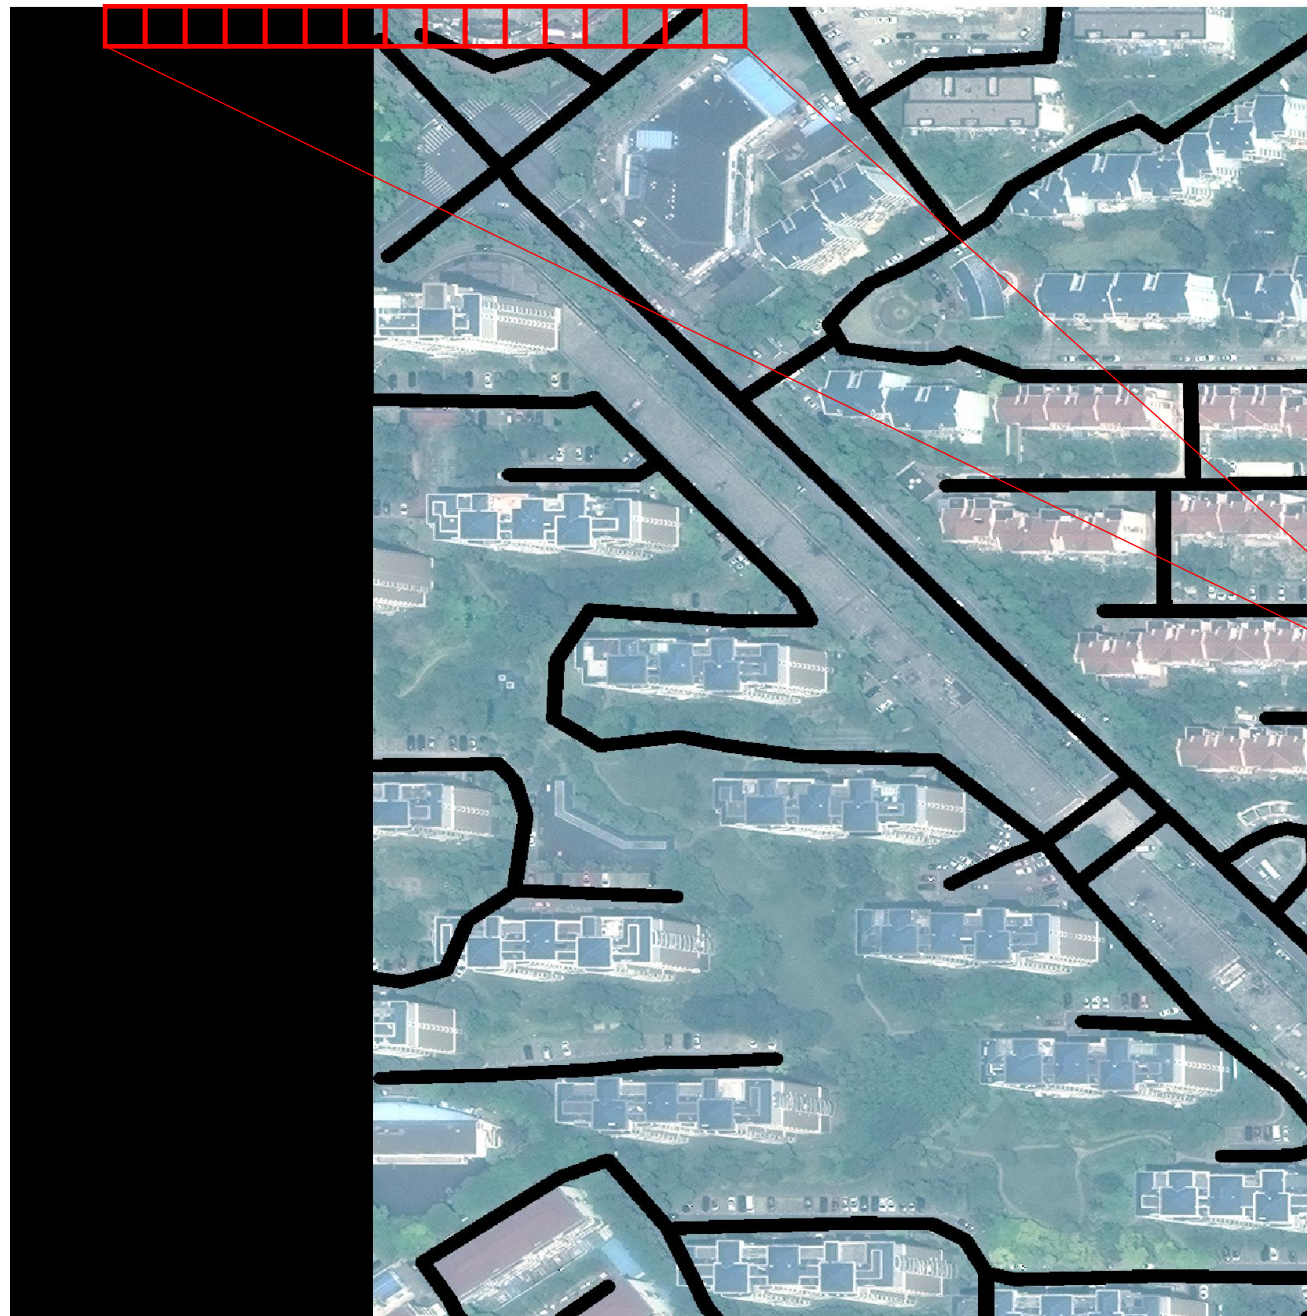


Hidden Layer

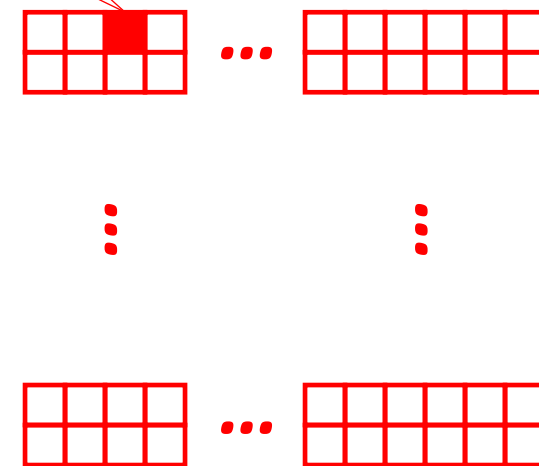


**A (16, 1) instead of (3, 3) Kernel (*not in real scale*)**





Hidden Layer



**A (16, 1) instead of (3, 3) Kernel (*not in real scale*)**

# Some Directions for the Next Step

Rotate the images for augmentation

Figure out a suitable kernel size

Create buffers by road type

Learn about transfer learning which uses pre-trained models

Read more papers on road segmentation

Try to train more images

